

Semantic MapReduce: Orchestrating LLMs over Large-Scale Data

Anonymous Author(s)

Abstract

Large-scale data systems efficiently scan, join, and aggregate massive datasets, but they do not directly solve semantic interpretation tasks such as evidence fusion, incident hypothesis generation, explanation, and auditable reasoning. Directly attaching an LLM to a database, log store, or dashboard is also insufficient: context is too large, intermediate evidence is hard to audit, and global conclusions must be derived from partitioned observations. This paper studies how an LLM-augmented system should orchestrate large-scale analysis when the hard part is not data access but semantic reduction. We frame the core abstraction as *Semantic MapReduce*: a small operator vocabulary for sharding observations, extracting local evidence, normalizing and grouping evidence objects, semantically reducing them into hypotheses, and generating traceable reports. We instantiate this abstraction in an anonymized orchestration-layer prototype that turns LLM reasoning workflows into explicit dataflow/stream pipelines, keeps model serving and data engines behind integration contracts, and makes evidence objects and workflow traces first-class artifacts. We introduce a reproducible large-scale analysis workload based on NPU-backed LLM serving telemetry. Across a 10-seed matrix and two workload sizes, a tail-aware deterministic incident reducer reaches mean precision 0.9500, mean recall 0.9750, and mean F1 0.9579, compared with 0.3075 mean F1 for a map-only alert baseline and 0.7129 for a window-aggregation baseline. An operator ablation shows that the missing incidents in the larger setting come from the map evidence policy rather than the reducer alone.

1 Introduction

Modern data systems are very good at moving, filtering, joining, and aggregating data. Hadoop MapReduce made large-scale analysis programmable through local map tasks followed by global reduction [2]; Spark and Flink improved the execution model for iterative, in-memory, batch, and stream workloads [1, 8]; Ray generalized distributed execution for emerging AI applications [6]. These systems are essential infrastructure, but they do not directly answer a different class of analysis questions: are these distributed symptoms part of the same incident, what evidence supports the hypothesis, which evidence conflicts with it, and what action should an operator take next?

This distinction appears even in a controlled telemetry workload. Figure 1 contrasts a naive stitched pipeline with the Semantic MapReduce design. In the NPU-backed LLM

serving workload, map-only alerts reach 0.3075 mean F1 and emit 3.50 outputs per true incident; window aggregation improves F1 to 0.7129 but still emits 1.76 fragments per incident. An incident-level reducer reaches 0.9579 F1 while reducing the report load to 1.04 outputs per true incident. The improvement therefore comes from changing the analysis contract itself: the system must turn shard-local observations into evidence objects and reduce them into incident hypotheses with traceable support, so reduction becomes a semantic operation rather than only a numeric or windowed aggregate.

Large language models (LLMs) make this semantic layer newly programmable. An LLM can summarize logs, classify symptoms, explain anomalous metrics, and produce operator-facing narratives. However, a simple “LLM over the database” architecture is a poor systems boundary. The input may contain millions of events, traces, metrics, documents, and experiment outputs. Relevant observations are partitioned by service, tenant, region, time window, and storage system. The output must be auditable: a conclusion should link back to concrete evidence objects rather than only to a transient prompt. Finally, cost, latency, privacy, and safety constraints make it undesirable to send all raw telemetry to a general-purpose model.

We therefore study the following systems problem: *how to orchestrate LLM-augmented large-scale analysis when the central operation is reducing distributed evidence into traceable semantic conclusions.*

We argue that the right abstraction is neither a monolithic agent nor a replacement for existing data engines. It preserves the local/global structure that made MapReduce effective, while changing the meaning of reduction. In traditional analytics, reduce computes values such as sums, counts, maxima, joins, and group-by aggregates. In LLM-native analysis, reduce combines evidence, deduplicates related symptoms, resolves conflicts, generates hypotheses, and produces explanations that remain linked to the observations that support them. We call this abstraction *Semantic MapReduce*, and define it around standard operators rather than around a single prompt or application script.

We instantiate this idea in an anonymized prototype. The prototype is a dataflow-native framework for modular, controllable, and transparent LLM-augmented reasoning. Its stream API exposes familiar dataflow composition primitives, its runtime provides local and optional distributed execution entrypoints, and its serving integration keeps inference engines outside the core package through explicit contracts.

Naive stitching





Figure 1. Motivating comparison. Naive stitching can send telemetry through scripts and an LLM agent, but the question of whether shard-local symptoms form one incident remains an implicit prompt convention. Semantic MapReduce makes the concrete boundary explicit: telemetry shards produce evidence objects, a semantic reducer merges them, and the report carries audit-trace links to the supporting evidence.

Rather than treating LLM analysis as a single prompt-response interaction, the prototype represents analysis as a workflow graph whose operators can be inspected, replayed, replaced, and connected to external systems. This makes the prototype an orchestration layer above or beside Spark, Flink, Ray, databases, vector systems, and LLM serving engines.

The resulting structure addresses cases that a chat interface over telemetry leaves implicit. The prototype can express a pipeline in which different shards are analyzed independently, local outputs are normalized into evidence objects, reduction logic is selected by policy, and final explanations remain connected to the intermediate decisions that produced them. The same workflow can run with a deterministic reducer, with a stub reducer for interface testing, or with an LLM-backed reducer for semantic fusion. The prototype therefore makes LLM-native analysis operationally structured.

The evidence boundary is deliberately narrow. The current workload does not implement production-grade distributed shuffle or exactly-once fault tolerance, and the main quality results do not claim LLM superiority over a deterministic reducer. It implements a reproducible workload, diagnostic baselines, a deterministic reducer, an `llm-stub` reducer for CI, and an OpenAI-compatible reducer path used for readiness probes. The supported claim is that the prototype expresses MapReduce-like semantic orchestration as

a standard operator workflow and exposes meaningful reducer failure modes. It is not a replacement for existing execution engines.

This paper makes five contributions:

1. It defines Semantic MapReduce as a systems abstraction and operator vocabulary for LLM-native analysis over partitioned data.
2. It maps the abstraction onto workflow, stream, runtime, evidence, and serving-integration boundaries, emphasizing how an orchestration layer can sit above existing data and execution engines.
3. It describes the mechanisms that make the abstraction practical: explicit workflow graphs, stream composition, reducer plug-ins, evidence objects, and auditable traces.
4. It introduces a reproducible large-scale analysis workload for NPU-backed LLM serving telemetry, including injected incidents, evidence operators, incident reduction, and precision/recall/F1 scoring.
5. It reports a 10-seed evaluation, an operator ablation, and real-endpoint readiness probes, showing where incident-level reduction helps, where map-stage evidence policy still determines recall, and why LLM reducers need an explicit JSON/evidence contract before quality claims are meaningful.

2 Motivation

What analytics engines leave to humans. Operational analysis often starts with numerical aggregates but ends with semantic judgment. Consider an LLM serving platform backed by accelerators. A dashboard may show that p95 latency increased in one region, decode utilization saturated, scheduler queue depth grew, and errors rose for some tenants. Each observation is measurable with conventional analytics. The harder question is whether these observations are one incident, several unrelated incidents, or noise. An operator must connect symptoms across services, identify the likely bottleneck, discount misleading signals, and decide what to inspect next.

Traditional execution engines provide the substrate for scans, windows, joins, and stateful aggregation. They do not normally provide a first-class abstraction for incident hypotheses, evidence objects, conflict handling, or explanation traces. These semantic tasks are often implemented as notebooks, scripts, alert rules, dashboards, and human reasoning. The result is powerful but hard to reproduce: two engineers may examine the same data and reach different conclusions, and the reasoning path may not survive the debugging session.

Why direct LLM integration is not enough. LLM-enabled data tools often expose a chat interface over a database, log store, vector index, or dashboard. This interface is useful for ad hoc exploration, but it hides important systems structure. The model must decide which context to inspect, how much evidence to compress, when to issue follow-up queries, and how to justify its answer. Without an explicit workflow, the analysis is difficult to schedule, parallelize, cache, replay, and audit.

Large-scale analysis needs a more structured design. First, raw data must be partitioned so each analysis unit fits within compute and context limits. Second, local outputs must be normalized into evidence objects with provenance. Third, reduction must be explicit, because it is the stage where duplicate evidence, contradictory signals, and weak hypotheses are accepted or rejected. Fourth, final reports must preserve traceability, especially if they are used in operational settings.

Why an orchestration layer is the right layer. The orchestration layer does not replace Spark, Flink, Ray, SQL engines, vector stores, or serving engines. Those systems already own important lower-level responsibilities: distributed execution, storage access, stream processing, indexing, model serving, and resource management. The orchestration layer instead expresses the LLM reasoning workflow as a dataflow/streaming pipeline, coordinates semantic operators, and records intermediate evidence. This layer is where sharding, evidence extraction, grouping, semantic reduction, and explanation generation become explicit.

3 Problem Statement

We define large-scale LLM-augmented data analysis as follows. The input is a large collection of heterogeneous observations: events, logs, metrics, traces, documents, experiment outputs, and derived aggregates. The output is a set of structured insights, incident hypotheses, explanations, evidence chains, and possible actions. A correct system must manage both scale and semantics.

This problem exposes six system challenges:

- **Scale and partitioning.** Data must be split across shards while preserving enough local context to extract meaningful evidence.
- **Evidence compression.** Raw records cannot all be sent to a reducer or LLM. Map stages must summarize without losing the signals needed for global reasoning.
- **Semantic reduction.** The reducer must merge related symptoms, remove duplicates, handle conflicts, and produce hypotheses rather than only numeric aggregates.
- **Auditability.** Reports must link back to the evidence and workflow decisions that produced them.
- **Latency and cost.** LLM calls, if used, must be bounded and placed where semantic value justifies cost.
- **Safety boundaries.** The orchestration layer must avoid leaking unnecessary raw data and must make it possible to restrict what reaches external models.

These constraints motivate a system that treats LLM interpretation as a workflow over evidence objects, not as a final chat step after data processing.

4 Semantic MapReduce

The central claim of Semantic MapReduce is that LLM-native analysis needs standard operators, analogous in spirit to map and reduce, but with contracts for evidence and explanation. We use six operators throughout the paper: SHARD partitions observations by scale, locality, policy, or context budget; MAPEVIDENCE extracts bounded local evidence; NORMALIZE converts local outputs into ranked evidence objects; GROUPEVIDENCE groups evidence by incident key, service, region, time, or semantic similarity; SEMANTICREDUCE merges, deduplicates, adjudicates conflicts, and emits hypotheses; and REPORTTRACE generates evidence-linked reports with operator provenance. External data engines may implement the scale-facing operators, while LLM frameworks may implement semantic reduction or reporting. The orchestration layer composes them under one auditable evidence contract.

Figure 2 summarizes the abstraction. The six operators form three logical stages.

Shard and map evidence. Each shard is processed locally. MAPEVIDENCE may compute statistics, identify anomalies, summarize logs, classify symptoms, or invoke an LLM over a bounded context. Its output is not arbitrary prose. It

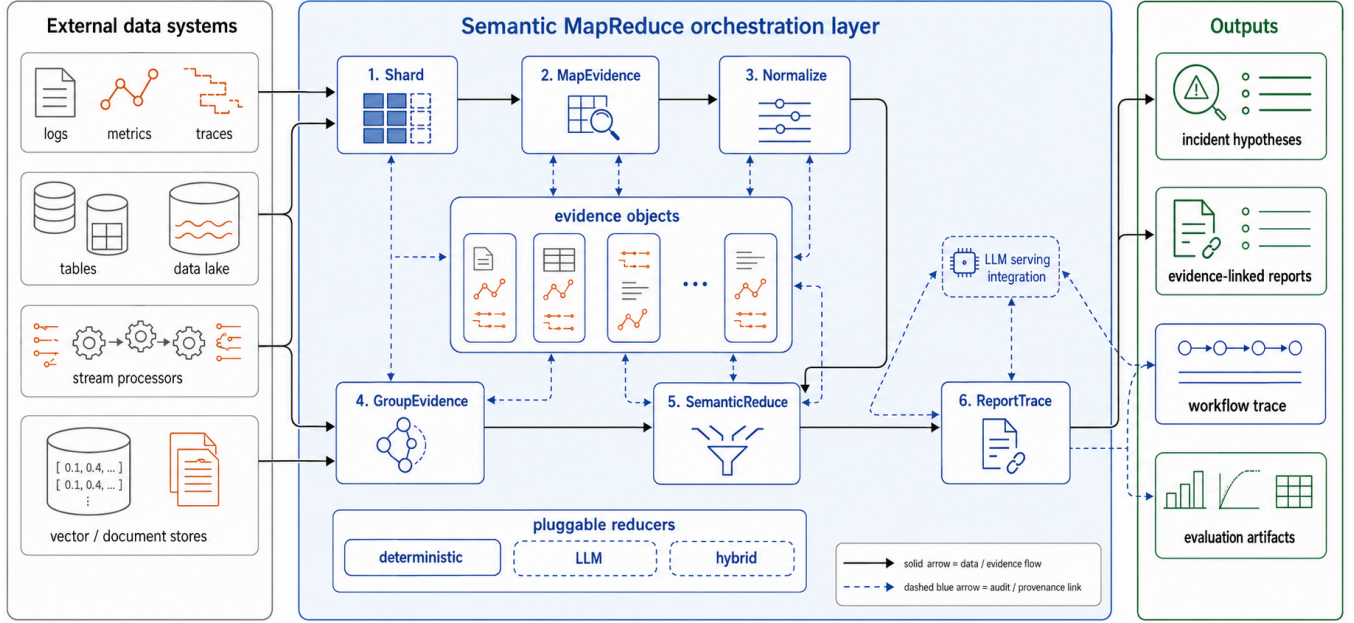


Figure 2. Semantic MapReduce as an orchestration layer over existing data and execution systems. External systems provide logs, metrics, traces, tables, streams, and vector or document stores. The orchestration layer coordinates SHARD, MAP EVIDENCE, NORMALIZE, GROUP EVIDENCE, SEMANTIC REDUCE, and REPORT TRACE around explicit evidence objects, with reducer and LLM-serving integration points kept behind pluggable boundaries. Outputs include incident hypotheses, evidence-linked reports, workflow traces, and evaluation artifacts.

is an evidence object with fields such as service, tenant, region, metric, time window, observed value, baseline, severity, confidence, and source references.

Normalize, group, and semantic reduce. NORMALIZE and GROUP EVIDENCE prepare evidence for global reasoning. SEMANTIC REDUCE then merges evidence objects into higher-level hypotheses. This stage deduplicates adjacent windows, combines weak but consistent signals, separates unrelated incidents, handles contradictory evidence, and assigns confidence. The reducer may be deterministic, LLM-backed, or hybrid. Deterministic reducers are useful baselines because they are reproducible and cheap; LLM reducers are attractive when the merge requires domain language, causal reasoning, or flexible explanation.

Report and trace. REPORT TRACE turns hypotheses into operator-facing reports. The report should include the claim, supporting evidence, possible conflicting evidence, and suggested actions. It should also preserve the workflow trace so a user can inspect how the conclusion was formed.

Semantic MapReduce differs from an agent loop because its units of parallelism, evidence compression, reduction, and reporting are explicit. It also differs from ordinary stream processing: stream windows can compute features and alerts, but semantic reduction turns local findings into a global explanation.

Algorithm 1 Semantic MapReduce Workflow

Require: Observations D , operators S, M_E, N, G_E, R_S, G_T

Ensure: Evidence-linked incident report $\mathcal{Y}$ and workflow trace $\mathcal{T}$

```

1:  $\mathcal{S} \leftarrow S(D)$  ▷ SHARD by service, region, tenant, or time
2:  $\mathcal{E} \leftarrow \emptyset, \mathcal{T} \leftarrow \emptyset$ 
3: for all shard  $s \in \mathcal{S}$  do
4:    $(E_s, T_s) \leftarrow M_E(s)$  ▷ MAP EVIDENCE
5:    $\mathcal{E} \leftarrow \mathcal{E} \cup E_s$ 
6:    $\mathcal{T} \leftarrow \mathcal{T} \cup T_s$ 
7: end for
8:  $\mathcal{E}' \leftarrow N(\mathcal{E})$  ▷ NORMALIZE
9:  $\mathcal{C} \leftarrow G_E(\mathcal{E}')$  ▷ GROUP EVIDENCE
10:  $(\mathcal{H}, T_R) \leftarrow R_S(\mathcal{C})$  ▷ SEMANTIC REDUCE
11:  $\mathcal{T} \leftarrow \mathcal{T} \cup T_R$ 
12:  $\mathcal{Y} \leftarrow G_T(\mathcal{H}, \mathcal{E}', \mathcal{T})$  ▷ REPORT TRACE
13: return  $(\mathcal{Y}, \mathcal{T})$ 

```

Algorithm 1 makes the operator vocabulary operational. The important constraint is that SEMANTIC REDUCE consumes compact, normalized, grouped evidence objects rather than raw telemetry. This is what lets the orchestration layer place expensive or policy-sensitive LLM calls at semantic boundaries instead of at every scan or filter operation.

5 System Design

5.1 Design Goals and Non-goals

The prototype targets four design goals for Semantic MapReduce.

Explicit workflow. Analysis is represented as a graph of operators so that task decomposition, intermediate evidence, and final reports are visible.

Composable streams. The system reuses stream/dataflow idioms for map, filter, flatmap, key grouping, and operator connection.

Pluggable semantic operators. The map and reduce stages support deterministic logic, LLM-backed logic, or hybrids under the same evidence schema.

Auditable evidence. The reducer produces structured incidents that reference the evidence used to form them.

The prototype also has explicit non-goals. It does not implement a full distributed shuffle for this workload. It does not provide exactly-once fault tolerance for this workload. It does not replace external data engines or model serving runtimes. These boundaries keep the current contribution focused on orchestration and workload design.

5.2 System Boundary and Workflow Model

Figure 2 also shows the intended system boundary. External systems provide raw observations, pre-aggregation, indexing, and model serving. The orchestration layer coordinates the semantic workflow: SHARD, MAP EVIDENCE, NORMALIZE, GROUP EVIDENCE, SEMANTIC REDUCE, and REPORT TRACE.

This boundary is important for comparison. Spark, Flink, and Ray can execute large computations; an orchestration layer can call or sit beside them when semantic orchestration is needed. LlamaIndex, LangGraph, LangChain, and AutoGen provide useful LLM application abstractions, but they are not primarily dataflow runtimes for auditable large-scale evidence reduction. Data+AI platforms provide deep integration, but this design emphasizes an open workflow/runtime boundary and explicit evidence trace.

Workflow graph as the operator IR. The workflow graph serves as an intermediate representation for LLM-augmented reasoning. A direct LLM integration often leaves the reasoning path inside a prompt chain or application glue code. In contrast, a workflow decomposes analysis into the standard operators introduced above. MAP EVIDENCE has a different role from NORMALIZE; SEMANTIC REDUCE has a different role from REPORT TRACE. This separation gives the system several practical properties.

First, the workflow can be inspected. A developer can see where evidence is extracted, where it is compressed, and where global hypotheses are formed. Second, the workflow can be replayed. If a reducer emits a wrong incident, the system can rerun the same map outputs through a new reducer instead of regenerating all raw telemetry. Third, operators can be replaced independently. A deterministic reducer can

be used for CI and regression tests, while an LLM-backed reducer can be evaluated under the same schema. Fourth, the workflow graph gives the runtime a place to record provenance: each final incident can be connected to the map outputs and reducer decision that produced it.

This is the sense in which the system is more than a prompt wrapper. The graph is not only a programming convenience; it is the unit of reproducibility and auditability for semantic analysis.

Stream composition as the data-system interface. The second advantage is that the system expresses LLM reasoning through dataflow and stream composition rather than through one-off callbacks. The stream API exposes operators such as map, filter, flatmap, key grouping, and stream connection. These primitives are familiar from data systems, but here they instantiate the Semantic MapReduce vocabulary: sharded map stages implement MAP EVIDENCE, keyed stages implement GROUP EVIDENCE, and downstream stages invoke SEMANTIC REDUCE and REPORT TRACE.

This stream boundary also gives the system a clean integration path. Upstream systems can provide already-partitioned events, pre-aggregated windows, or document chunks. The orchestration layer does not need to own the storage engine to coordinate the semantic workflow. Downstream systems can consume incident objects, reports, or traces. The system is therefore complementary: it adds a semantic orchestration plane without owning every execution primitive below it.

Evidence objects as the semantic data model. The workload uses structured evidence rather than free-form text. A simplified evidence object contains:

- shard and time-window identifiers;
- service, region, tenant, and metric fields;
- observed values and local baseline statistics;
- anomaly type and severity;
- confidence and source-event references.

This schema is intentionally model-agnostic. A deterministic mapper can populate it from thresholds and percentiles. A future LLM mapper can populate it from logs or traces. A reducer can then operate over the same evidence representation. The key design choice is that evidence is not merely embedded in prose. It is a first-class object that can be scored, matched to ground truth, serialized in reports, and inspected after the workflow completes.

Evidence objects are the bridge between data systems and LLM reasoning. Data engines usually produce rows, windows, and aggregates. LLMs often produce text. Semantic MapReduce needs an intermediate form that is structured enough for systems evaluation and expressive enough for semantic interpretation. The evidence object plays that role.

Table 1. Division of responsibility in the Semantic MapReduce stack. The prototype is positioned as the orchestration layer rather than a replacement for data engines or model-serving systems.

Layer	Primary responsibility	Why it matters for LLM-native analysis
Data engines	Storage access, scans, joins, stream windows, pre-aggregation, distributed execution	Keeps high-volume data movement in systems already optimized for it; avoids making the orchestration layer a replacement for Spark, Flink, Ray, or databases.
LLM serving engines	Model execution, batching, scheduling, tokenizer/model-specific serving policy	Keeps model serving modular; a semantic reducer can call external models without embedding a serving engine in the workflow layer.
Workflow/stream run-time	Operator graph, stream composition, map/reduce orchestration, reducer selection, trace capture	Makes LLM reasoning explicit, replayable, and inspectable instead of hiding it in a prompt chain.
Semantic operators	Evidence extraction, hypothesis merging, conflict handling, explanation generation	Encodes the semantic work that conventional aggregations do not express: merging symptoms into auditable conclusions.
Evaluation harness	Ground truth, precision/recall/F1, missed incidents, reducer comparisons	Provides a reproducible way to compare deterministic, LLM-backed, and hybrid reducers on identical workloads.

5.3 Reducers, Traces, and Integration

The current code introduces a reducer abstraction with four concrete roles: a map-only diagnostic baseline, a window-aggregation baseline, a deterministic incident reducer, and an llm-stub reducer. The map-only baseline reports shard-local candidates directly. The window-aggregation baseline merges candidates within each service/region/window but does not cluster adjacent windows into incident-level hypotheses. The deterministic reducer performs that incident-level clustering using structured fields. The stub reducer preserves the future LLM interface while delegating to the deterministic baseline so CI and offline experiments remain reproducible.

The reducer interface is the extension point for LLM-backed reduction. An LLM reducer can take evidence objects, call a model under bounded context, and emit the same incident schema. Because scoring is unchanged, deterministic and LLM reducers can be compared on identical seeds and ground truth. This is an important systems property: changing the semantic reducer does not require rewriting the generator, shard mapper, scorer, report schema, or benchmark matrix.

Algorithm 2 shows the reducer interface used to separate evidence production from hypothesis formation. The current implementation includes diagnostic non-semantic reducers as well as the deterministic incident reducer; the llm-stub branch preserves the LLM call boundary. An LLM implementation changes how evidence is merged into hypotheses while keeping the incident schema and scorer fixed.

Workflow trace and auditability. The trace is the mechanism that makes semantic conclusions auditable. In an incident-analysis workflow, the final answer is only useful if a human can inspect why it was produced. The trace records

Algorithm 2 Reducer Contract Used by the Workload

Require: Ranked evidence objects $\mathcal{E}$, reducer policy π , evidence budget k
Ensure: Incident hypotheses $\mathcal{H}$ and reducer trace T_R

```

1:  $C \leftarrow \text{GROUPCANDIDATES}(\mathcal{E}, \text{type}, \text{service}, \text{region}, \text{window})$ 
2:  $\mathcal{H} \leftarrow \emptyset, T_R \leftarrow \emptyset$ 
3: for all candidate group  $c \in C$  do
4:    $E_c \leftarrow \text{SELECTTOPEVIDENCE}(c, k)$ 
5:   if  $\pi = \text{DETERMINISTIC}$  then
6:      $h \leftarrow \text{RULEMERGE}(E_c)$ 
7:   else if  $\pi = \text{LLM}$  then
8:      $h \leftarrow \text{LLMMERGE}(E_c)$  ▷ Future implementation
9:   else
10:     $h \leftarrow \text{HYBRIDMERGE}(E_c)$  ▷ Rules filter, LLM adjudicates
11:   end if
12:   if  $\text{ISVALIDINCIDENT}(h)$  then
13:      $\mathcal{H} \leftarrow \mathcal{H} \cup \{h\}$ 
14:      $T_R \leftarrow T_R \cup \text{TRACEDECISION}(h, E_c, \pi)$ 
15:   end if
16: end for
17: return  $(\mathcal{H}, T_R)$ 

```

which shard summaries were considered, which evidence objects were merged, which incident hypothesis each evidence object contributed to, and which injected incidents were missed during evaluation. The current workload records detected incidents, injected incidents, matched incident IDs, and missed incident IDs; the same trace can later be exposed through an interactive visualization layer.

Auditability matters for both correctness and operations. A false positive is explainable as a specific cluster of noisy evidence rather than as an opaque model hallucination. A false negative is linked to missing or weak evidence rather

than only to a low aggregate score. The workflow orientation gives the system a natural place to attach trace information at each operator boundary.

External integration points. The prototype can integrate with external systems at three points. Upstream, data engines can supply shards, windows, or materialized feature tables. In the middle, model-serving systems can provide LLM calls for map, reduce, or final explanation operators. Downstream, incident-management or observability tools can consume structured hypotheses and evidence-linked reports. The current workload uses a synthetic generator and deterministic reducer, but the interfaces are chosen so that real Spark/Ray/Flink scans, real vector retrieval, and real serving telemetry can be substituted later.

This integration path is central to the design. The system does not duplicate mature execution engines. It makes semantic orchestration explicit across systems that otherwise remain disconnected: databases, stream processors, vector stores, LLM serving endpoints, and human-facing incident reports.

Control plane and data plane. Another way to describe the boundary is to separate the semantic control plane from the data plane. The data plane moves records, executes scans, stores logs, manages model kernels, and serves tokens. These performance-critical tasks remain in engines designed for them. The semantic control plane decides how analysis is decomposed, which summaries become evidence, which reducer combines that evidence, and how the final report exposes provenance. The prototype belongs primarily to this control plane.

This distinction is useful because LLM-native analysis often fails when the control plane is hidden inside an application script. For example, a notebook can query a database, call an LLM, and write a report, but the boundaries between data extraction, evidence selection, hypothesis formation, and explanation are implicit. Reproducing the result requires reproducing the notebook state and the human decisions around it. In our prototype, these steps are represented as workflow operators and report fields. The system can therefore record the exact reducer used, the top-k evidence budget, the seed, the matched incidents, and the missed incidents.

The control-plane view also clarifies the integration boundary. A distributed deployment may push map operators into Spark or Ray, stream events through Flink, or call an LLM endpoint for semantic reduction. In each case, the orchestration layer owns the orchestration contract: operator graph, reducer selection, evidence schema, trace capture, and report generation. This is the level at which the contribution is strongest.

Why the design is AI-native. In this workload, the prototype is AI-native in a concrete sense: the system interface

treats semantic operators as first-class components. A reducer is not only a function from rows to rows; it is a component that may reason over evidence, produce a hypothesis, cite support, expose uncertainty, and preserve a trace for later inspection. The workflow is not only an execution plan; it is also a reasoning plan.

This design has consequences for implementation. The report schema carries evidence references, not only output strings. The reducer interface is pluggable because deterministic, LLM-backed, and hybrid reducers have different cost and reliability profiles. The benchmark records false positives and false negatives in inspectable form because aggregate F1 alone does not tell an operator whether a missed incident was caused by weak evidence, poor clustering, or an overaggressive threshold. These requirements arise specifically because the workload includes LLM interpretation.

6 Workload

6.1 Scenario

The workload models telemetry from an NPU-backed LLM serving platform. Services include `prefill`, `decode`, `kv-cache`, `scheduler`, `router`, and `embedding`. Events include `service`, `tenant`, `region`, `timestamp`, `latency`, `error flag`, `NPU utilization`, `queue depth`, and `generated metadata`. The generator injects hidden incidents of three types:

- **Latency spike:** one service experiences elevated latency over a time window.
- **NPU saturation:** utilization increases and may affect queueing or latency.
- **Queue backlog:** scheduler or routing queues grow and create downstream symptoms.

The task is to recover injected incidents from event streams. A detected incident is useful only if it links to the supporting evidence and matches the hidden ground truth by type, service, region, and time.

Why this workload exercises the system. The workload is designed to stress the orchestration layer rather than raw storage throughput. It has enough structure to require semantic reasoning: incidents are injected across services, regions, and time windows, while local symptoms appear as latency, error, utilization, or queue-depth changes. A single shard can observe symptoms, but the global conclusion depends on how those symptoms are merged. This matches the Semantic MapReduce pattern.

The workload also forces the system to preserve evidence. If the reducer reports a latency spike, the report must include which shard-level candidates supported the decision. If it misses a saturation incident, the report should expose the missed incident ID. This is why the workload reports both detected and injected incidents, not only aggregate F1. The goal is not merely to score a classifier; it is to evaluate an auditable analysis workflow.

Finally, the workload is intentionally reproducible. Seeds, event counts, shard counts, top-k values, reducer names, and incident reports are recorded. The same matrix can be rerun when a reducer changes. This gives the prototype a benchmark harness for future LLM reducer experiments without making the current paper depend on a live model endpoint.

6.2 Pipeline and Implementation

The workload instantiates the six operators directly. SHARD partitions generated telemetry by shard count, service, region, and time. MAPEVIDENCE computes shard-level statistics, anomaly candidates, and local summaries. NORMALIZE converts candidates into scored evidence objects with provenance fields. GROUPEVIDENCE clusters candidates by incident type, service, region, and temporal overlap. SEMANTICREDUCE emits incident hypotheses with map-only, window, deterministic, or future LLM reducers. REPORTTRACE scores hypotheses, records matched/missed incidents, and emits per-run reports.

The workload reports precision, recall, F1, throughput, map latency, reduce latency, detected incidents, injected incidents, and missed incidents. Reports also include seed, top-k, reducer name, and incident matching metadata.

Current implementation. The current implementation has three pieces. The generator creates synthetic telemetry and hidden incidents. The map stage implements SHARD, MAPEVIDENCE, and NORMALIZE by partitioning events, computing local summaries, and producing anomaly candidates. The reducer implements GROUPEVIDENCE and SEMANTICREDUCE over shard summaries. The default reducer is deterministic. It clusters candidates using incident type, service, region, and temporal overlap. The map-only and window-aggregation reducers are diagnostic baselines that approximate alert-only and windowed-analytics workflows without incident-level semantic reduction. The `llm-stub` reducer is an interface placeholder that exercises the same resolver and report path but delegates to the deterministic reducer.

The report schema is deliberately richer than the minimum needed for precision and recall. It includes the map policy, reducer name, seed, top-k, injected incidents, detected incidents, matched incident IDs, and missed incidents. This makes the artifact useful for debugging operator behavior. When a run misses an incident, the missed incident ID is preserved in the report rather than disappearing into the aggregate recall number; when a map policy changes, the same reducer and scorer can be replayed over the new evidence stream.

Case study walkthrough. A typical run proceeds as follows. The generator emits events for services such as `prefill`, `decode`, and `scheduler`. It injects incidents such as a latency spike or NPU saturation over a bounded time window. The shard map stage computes local evidence: elevated p95 latency, higher queue depth, or high NPU utilization. Each

Table 2. Workload configurations. Each configuration is run with ten seeds (7, 11, 13, 17, 19, 23, 29, 31, 37, and 41), two map policies, and four reducer regimes.

Events	Shards	Top-k	Runs
50,000	16	12	80
100,000	32	16	80

local summary is compact enough to be passed to the reducer. The reducer then merges adjacent or related candidates into incident hypotheses. The final report presents incident type, service, region, time window, confidence-like scores, and matching information.

This walkthrough is simple, but it illustrates the value proposition. The workload is not a single SQL query and not a single prompt. It is a workflow whose operators correspond to system responsibilities: sharding, local evidence extraction, normalization, grouping, semantic merging, reporting, and evaluation. Each responsibility can evolve independently. A future LLM reducer can replace the deterministic reducer without changing the generator or scorer; a real telemetry source can replace the synthetic generator without changing the reducer contract.

Experimental provenance. The workload is designed to make reducer behavior reproducible without placing environment-management details in the paper body. Each experiment records the workload scale, shard count, seed, reducer, top-k budget, detected incidents, missed incidents, and latency measurements. This provenance is sufficient to distinguish algorithmic changes in the reducer from changes in input generation or scoring. The accompanying artifact contains the implementation and per-run reports needed to rerun the study.

7 Evaluation

7.1 Setup and Scope

Table 2 lists the reproduced configurations. The evaluation is deliberately scoped to operator behavior under one workload harness. It does not claim tuned cluster deployments of Spark, Flink, Ray, LangGraph, LlamaIndex, or AutoGen. The main comparison fixes the generator, evidence schema, scorer, and report format, then varies two pieces: the map evidence policy and the reducer. This lets us separate three questions that were conflated in earlier drafts: whether incident-level reduction helps, whether the map stage preserves enough evidence for any reducer to succeed, and whether adjacent orchestration frameworks provide the missing semantic contract by themselves.

7.2 Reducer Quality

The aggregate results in Table 3 show why incident-level reduction is the relevant operation. Map-only alerts emit




Figure 3. Evaluation harness for the large-scale analysis workload. The same generated telemetry and injected incidents feed the evidence stage, reducer variants, and common scorer. Solid reducer rows correspond to implemented comparison regimes, while the dashed LLM/hybrid branch marks planned extensions rather than completed LLM-reducer results. Ground truth is used only by the scorer, and workflow-trace artifacts link results back to evidence objects.

Table 3. Reducer comparison with the tail-aware map policy across 20 runs. Map-only and window-aggregate are diagnostic baselines for alert-style and windowed-analytics workflows. The llm-stub row delegates to the deterministic reducer and is included only as an interface check.

Reducer	Precision	Recall	F1	Detections
Map-only	0.1990	0.6875	0.3075	14.00
Window-aggregate	0.5696	0.9750	0.7129	7.05
Incident reducer	0.9500	0.9750	0.9579	4.15
llm-stub	0.9500	0.9750	0.9579	4.15

many duplicate shard-local candidates and consume the top-k budget before all incidents are represented. Window aggregation improves recall but still reports multiple windows for one incident. The incident reducer keeps recall at 0.9750 while reducing average detections from 7.05 to 4.15, raising F1 from 0.7129 to 0.9579. This is the measured advantage of semantic reduction in the current workload: fewer duplicate hypotheses, not faster scanning.

Table 4 is the most important diagnostic result. The larger workload exposes a failure that is not solved by changing the reducer alone: with mean-only NPU detection, three of ten 100K runs miss at least one incident. The missed cases include NPU saturation windows whose mean utilization is diluted by surrounding normal events, even though their tail utilization is high. Adding a stricter p95 NPU signal to MAP EVIDENCE reduces missed runs from three to one and

Table 4. MapEvidence policy ablation for the deterministic incident reducer. The tail-aware policy preserves high-p95 NPU saturation evidence; mean-only corresponds to the earlier mean-utilization rule.

Size	Map policy	Precision	Recall	F1	Missed runs
50K/16/12	Mean-only	0.9200	0.9750	0.9413	1/10
50K/16/12	Tail-aware	0.9200	0.9750	0.9413	1/10
100K/32/16	Mean-only	0.9750	0.9250	0.9464	3/10
100K/32/16	Tail-aware	0.9800	0.9750	0.9746	1/10

raises 100K F1 from 0.9464 to 0.9746. This result strengthens the operator argument: once map, evidence, and reduce are separate contracts, a failure can be localized to the evidence policy rather than vaguely attributed to “the LLM” or “the reducer.”

The ablation also prevents overclaiming. Tail-aware evidence is not a universal improvement: at 50K it is neutral, and an overly permissive p95 threshold creates false positives. The reported threshold is therefore a tuned policy for this workload, not a general saturation detector. A stronger system would learn or calibrate the evidence policy from real traces.

7.3 Runtime and Adapter Overhead

The local runtime metrics in Table 5 show that the map stage dominates measured workload time because it scans and summarizes event shards. The reduce stage is sub-millisecond

Table 5. Local runtime metrics for the tail-aware deterministic reducer, averaged over ten seeds. These numbers support reproducibility and regression tracking; they are not cluster-scale performance results.

Configuration	50K/16/12	100K/32/16
Throughput, events/s	105,215	102,819
Map latency, ms	73.50	162.17
Reduce latency, ms	0.30	0.29
Total latency, ms	475.34	972.91

Table 6. Adjacent-framework adapter baselines over the same 20 tail-aware runs. Each adapter receives the same evidence objects and calls the same incident reducer and scorer; equal quality is therefore expected.

Adapter	F1	Throughput	Total ms
Native-local	0.9579	58,136	1,303.22
LangGraph-local	0.9579	58,184	1,299.13
LlamaIndex-docstore	0.9579	57,868	1,304.96
Ray-local	0.9579	39,100	1,917.48

in the current baseline because it performs deterministic clustering over compact candidate summaries. The workload therefore exercises the intended architecture: raw event processing precedes semantic reduction, and the reducer sees compressed evidence rather than all raw events.

These numbers are not distributed systems performance claims. The current runs are local and do not evaluate cluster scheduling, shuffle, or fault recovery. They establish a reproducible baseline for later reducer studies. If an LLM reducer improves recall but adds latency or token cost, the same harness can quantify the tradeoff. A complete evaluation would report end-to-end latency and cost for deterministic, LLM-backed, and hybrid reducers under the same workload.

Adjacent-system baselines. Table 6 adds adjacent-framework baselines for Ray-local, LangGraph-local, and LlamaIndex-docstore wrappers. The table should not be read as a leaderboard over those systems: the experiment intentionally fixes the evidence schema, incident reducer, scorer, and report format. Its value is diagnostic. Without semantic reduction, the workflows in Table 3 stop at shard-local alerts or window fragments. Once the semantic-reduce contract is supplied, adjacent frameworks can drive the same computation and obtain identical quality. This supports a narrower claim: existing systems are useful substrates, but the missing abstraction for this workload is the evidence-linked reducer contract and its audit trail, not another generic execution wrapper. Ray-local also reported NPU-detection and /tmp/ray space warnings on this host, so its overhead should not be read as tuned Ray performance.

Table 7. Single-NPU online endpoint run. The endpoint served Qwen2.5-7B on one Ascend 910B2 NPU. Each row reports eight measured streaming requests after warmup. The table validates endpoint integration and measurement plumbing rather than SOTA serving performance.

	Concurrency	OK	TTFT ms	TPOT ms	Tok/s
1		8/8	117.69	70.18	13.82
2		8/8	1,494.64	70.24	13.883

7.4 Real-Online Readiness

The online endpoint results in Table 7 attach a live model-serving endpoint to the workload effort. The purpose is not to benchmark serving throughput, but to validate that the reducer path can be connected to an external LLM service and measured with standard serving metrics. The concurrency-2 row has much larger TTFT because requests queue behind the active decode in this constrained setup. The similar TPOT values indicate that the decode path itself is stable across the two small runs.

Endpoint health is not sufficient for a semantic reducer. Follow-up online probes against an NPU3 Qwen2.5-7B endpoint were able to launch the service, pass health checks, and execute authenticated chat requests after local vLLM-HUST compatibility fixes, but the model did not satisfy the reducer JSON contract: an incident-shaped prompt produced repeated non-JSON text, structured output entered the schema path but did not close valid JSON, and even a minimal JSON prompt produced repeated non-JSON text. We therefore treat these probes as serving/generation readiness evidence, not as LLM reducer quality results. The lesson for the system design is that SEMANTICREDUCE needs a first-class contract and readiness gate; otherwise a nominally live LLM endpoint can silently invalidate an experiment.

Audit artifacts. The benchmark produces artifacts beyond a single summary table. Each run records the configuration, reducer name, injected incidents, detected incidents, matched incident IDs, missed incident IDs, and latency measurements. This matters because reducer experiments require more than a single score. A researcher should be able to identify which incident changed and inspect the evidence path that caused the change.

The artifact structure also supports regression testing. If a code change improves mean F1 but introduces a systematic false positive in one service or region, the per-run reports expose that pattern. If an LLM reducer improves recall only by emitting many vague incidents, precision falls and the detected incident list shows why. If an LLM reducer is too expensive, latency and token accounting can be attached to the same report schema. In this sense, the workload is not only a synthetic generator; it is an experimental harness for studying semantic reducers under controlled conditions.

7.5 Lessons

The evaluation gives three lessons. First, the workload is not merely a classifier benchmark. Precision measures whether the reducer avoids inventing incidents from noisy evidence; recall measures whether the map and reduce stages preserve weak but real incidents; the missed-incident list shows which operator boundary failed. Second, the incident reducer improves quality mainly by changing the unit of output from windows to hypotheses. The output count drops toward the four injected incidents while recall remains high. Third, the map policy matters as much as the reducer: if a saturation event never becomes evidence, no downstream LLM can recover it without going back to raw telemetry.

These lessons define the role of an LLM reducer more narrowly than the introduction’s motivating vision. A model should not be rewarded for basic scanning, thresholding, or duplicate removal that deterministic operators already handle. Its useful role is to adjudicate ambiguous evidence groups, use service relationships to reject false positives, explain why symptoms belong together, and produce operator-facing reports from bounded evidence objects. That experiment remains future work; the current artifact establishes the operator contract and the baseline failure modes that such a reducer must beat.

8 Discussion

The prototype is best viewed as a semantic control plane, not a replacement data plane. Data engines still own storage access, scans, joins, stream windows, and distributed execution. Model-serving engines still own inference. The orchestration layer owns a different contract: which observations become evidence, which reducer policy is applied, how hypotheses cite evidence, and how failures are traced. This is why the paper emphasizes operators rather than a single agent loop.

The current results also delimit the contribution. They show that the workflow boundary can run a nontrivial Semantic MapReduce workload end to end and make operator failures visible. They do not show production distributed execution, exactly-once stream semantics, or LLM superiority over a deterministic reducer. The next system step is a hybrid reducer that consumes the same evidence objects, compares deterministic and LLM decisions under the same scorer, and records token cost, model latency, and evidence coverage.

9 Related Work

Large-scale data processing. MapReduce, Spark, and Flink provide the execution substrate for large-scale batch and stream analytics [1, 2, 8]. Semantic MapReduce borrows the local/global decomposition but changes the reducer contract from fixed aggregation to semantic evidence fusion.

Distributed AI execution. Ray provides a general distributed runtime for AI and Python applications [6]. The prototype can use such systems as execution substrates, but its focus is the explicit orchestration of LLM reasoning workflows and evidence traces.

LLM serving. vLLM shows how memory management and scheduling shape LLM serving performance [3]. Our workload is motivated by serving telemetry in accelerator-backed deployments. This paper does not modify the serving engine; serving systems act as external sources of telemetry and possible model endpoints for semantic reducers.

LLM application frameworks. LangGraph, LlamaIndex, and AutoGen provide useful abstractions for agents, retrieval, tool use, and multi-agent workflows [4, 5, 7]. A full comparison requires verifying their dataflow, streaming, runtime, and audit semantics under the same workload. The narrower comparison here is that the prototype emphasizes workflow/stream/runtime orchestration with explicit semantic reduction and evidence provenance.

Data+AI platforms. Systems such as Databricks, Snowflake Cortex, and BigQuery ML integrate analytics engines with machine learning and LLM functionality. These platforms are important reference points, but they require systematic comparison before strong claims about relative capability. The prototype differs in emphasis: it is an open workflow/runtime layer that can sit above multiple engines and make reducer behavior auditable. The main comparison dimensions are APIs, execution boundaries, governance mechanisms, and traceability features.

Observability and incident management. Production observability systems collect metrics, logs, traces, alerts, and incident timelines. They provide the raw material for the workload studied here. The prototype does not replace them. Instead, Semantic MapReduce asks how evidence from these systems can be converted into auditable incident hypotheses. This is an important distinction: an alert can say that a threshold fired, while a semantic reducer should explain how multiple alerts and metrics combine into one operational hypothesis.

Agentic data analysis. Recent agentic systems can write queries, call tools, inspect results, and refine an answer. Such systems are useful for exploratory analysis, but their control flow may be difficult to reproduce when applied to high-volume operational data. The prototype chooses a more constrained structure for this paper: explicit map stages, explicit reducer stages, and explicit evidence objects. This reduces flexibility, but it improves replayability, evaluation, and audibility.

Positioning summary. The closest intellectual ancestor of Semantic MapReduce is MapReduce itself: local processing followed by global reduction. The closest practical neighbors are stream processors, distributed AI runtimes, LLM workflow frameworks, RAG/data frameworks, and data+AI

platforms. The prototype sits between these categories. It is not the storage engine, not the model server, and not merely an agent framework. It is the orchestration layer that makes semantic data analysis explicit enough to evaluate.

This positioning is intentionally cautious. The paper does not claim that the prototype is more scalable than Spark or Ray, more feature-complete than data+AI platforms, or more general than agent frameworks. The sharper claim is that the prototype gives this workload the right system structure: standard evidence operators, semantic reduction, evidence-linked reporting, and workflow traces under one reproducible contract.

10 Limitations

The current study has three main limitations. First, the workload is synthetic: it is modeled after accelerator-backed LLM serving telemetry, but it is not yet a replay of production traces. Second, the main evaluated semantic reducer is deterministic; the online endpoint probes validate integration and readiness checking, but they do not show that LLM reduction improves incident recovery over the deterministic baseline. Third, the experiments are local and do not evaluate distributed shuffle, exactly-once processing, fault recovery, or full end-to-end deployments of external systems.

These boundaries define the claim. The results support Semantic MapReduce as an orchestration abstraction and show that the prototype can coordinate shard-level analysis, evidence objects, incident reduction, and workflow traces under a reproducible contract. They do not show that the prototype replaces data engines or model-serving systems, nor that LLM reduction already improves incident recovery. The next step is to evaluate deterministic, LLM-backed, and hybrid reducers on real or replayed telemetry under the same evidence and scoring interface.

11 Conclusion

Large-scale LLM-augmented analysis needs more than fast scans and more than chat over data. It needs a workflow layer that can partition observations, compress evidence, semantically reduce hypotheses, and preserve audit trails. Semantic MapReduce captures this need by extending the map/reduce structure from numeric aggregation to evidence fusion and explanation. The prototype’s dataflow, stream, runtime, and serving boundaries make it a concrete substrate for this direction. The current workload demonstrates shard-level analysis and incident reduction while identifying the next experimental step: LLM-backed and hybrid reducers evaluated against the same evidence and scoring contract.

References

- [1] Paris Carbone, Asterios Katsifodimos, Stephan Ewen, Volker Markl, Seif Haridi, and Kostas Tzoumas. 2015. Apache Flink: Stream and Batch Processing in a Single Engine. *IEEE Data Engineering Bulletin* (2015).
- [2] Jeffrey Dean and Sanjay Ghemawat. 2004. MapReduce: Simplified Data Processing on Large Clusters. In *Proceedings of the 6th Symposium on Operating Systems Design and Implementation*.
- [3] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient Memory Management for Large Language Model Serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles*.
- [4] LangChain. 2026. LangGraph: Build Stateful, Multi-Actor Applications with LLMs. <https://github.com/langchain-ai/langgraph>. Accessed 2026-06-30.
- [5] LlamaIndex. 2026. LlamaIndex: Data Framework for LLM Applications. https://github.com/run-llama/llama_index. Accessed 2026-06-30.
- [6] Philipp Moritz, Robert Nishihara, Stephanie Wang, Alexey Tumanov, Richard Liaw, Eric Liang, Melih Elibol, Zongheng Yang, William Paul, Michael I. Jordan, and Ion Stoica. 2018. Ray: A Distributed Framework for Emerging AI Applications. In *Proceedings of the 13th USENIX Symposium on Operating Systems Design and Implementation*.
- [7] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, Ahmed H. Awadallah, Adam White, Doug Burger, and Chi Wang. 2023. AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation. [arXiv:2308.08155](https://arxiv.org/abs/2308.08155) [cs.AI]
- [8] Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, and Ion Stoica. 2012. Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing. In *Proceedings of the 9th USENIX Symposium on Networked Systems Design and Implementation*.